Autoregression and Moving Average

- Box-Jenkins (1976)
- No economic theory. For fitting and prediction only.

Time Series Probability Model

- $X(t, \omega)$
- For a fixed $t = 1, \ldots, T$, $X(t, \cdot)$ is a random variable
- For a fixed $\omega \in \Omega$, $X(\cdot, \omega)$ is a sequence

Ergodic Theorem

If the time series $X_t(\omega)$ is strictly stationary, and $E_t[X(\omega)] = \mu$, then as $n \to \infty$,

$$n^{-1} \sum_{t=1}^{n} X_t(\omega) \xrightarrow{p} \mu,$$

More formally,

$$P\left[\omega \in \Omega : \left|\frac{1}{n} \sum_{t=1}^{n} X_t(\omega) - \mu\right| > \epsilon\right] \to 0.$$

Temporary average $n^{-1} \sum_{t=1}^{n} X_t$ converges to spatial average $E[X_t(\cdot)]$.

A Simple Models

- White noise: $(e_t)_{t=-\infty}^{\infty}$:
    - $E[e_t] = 0$, $E[e_t^2] = \sigma_e^2$, and $E[e_t, e_s] = 0$ for all $t \neq s$.

ARMA

- AR(p)

$$y_t = \mu + \gamma_1 y_{t-1} + \cdots \gamma_p y_{t-p} + e_t$$

- MA(q)

$$y_t = \mu + e_t - \theta_1 e_{t-1} - \theta_q e_{t-q} + e_t$$

- ARMA(p,q)

$$(1 - \Gamma(L))y_t = \mu + \Theta(L)e_t$$

Stationarity: in AR form whether all roots lies out of the unit cycle.

Autocorrelation Pattern

- MA(q): finite dependence
- AR(1): geometric decline
    - $E[y_t] = \mu/(1 - \gamma_1)$
    - $\mathrm{var}[y_t] = \sigma_e^2/(1 - \gamma_1^2)$
    - $E[y_t|y_{t-1}] = \mu + \gamma_1 y_{t-1}$
    - $\mathrm{var}[y_t|y_{t-1}] = \sigma_e^2$

Modeling

- ARIMA(p, r, q)

$$(1 - \Gamma(L))\Delta^r y_t = \mu + \Theta(L)e_t$$

- Transform into stationary time series by taking logarithm and/or difference.
- Fit ARMA(p,q)

Seasonality

- Generated due to sampling frequency
- Add dummies to control seasonality

Estimation

- MLE for MA(q)
- MLE for ARMA(p,q)
- OLS for AR(p)

Model Choice

Information criteria.

Let k be the total number of slope coefficient in the model.

- Akaike information criterion: $\log(\hat{\sigma}^2) + 2 \times (k/T)$.
    - Tend to overfit, but better for prediction
- Bayesian information criterion: $\log(\hat{\sigma}^2) + \log(T) \times (k/T)$
    - Model selection consistent

Information criteria are not restricted to time series regressions. They are general statistical measures for model/variable selection.

AR(1) without drift

$$y_t = \gamma_1 y_{t-1} + e_t$$

$$\sqrt{T}(\hat{\gamma}_1 - \gamma_1) = \sqrt{T}\frac{\sum_{t=2}^{T} y_{t-1}e_t}{\sum_{t=2}^{T} y_{t-1}^2}$$

$$\xrightarrow{d} N\left(0, \frac{\operatorname{var}[y_t]\operatorname{var}[e_t]}{(\operatorname{var}[y_t])^2}\right)$$

$$\sim N\left(0, \frac{\sigma_e^2}{\sigma_e^2/(1-\gamma_1^2)}\right)$$

$$\sim N\left(0, 1 - \gamma_1^2\right)$$

What happens when $|\gamma_1| \to 1$?

Unit Root

- AR(1) with AR coefficient equals 1.

$$y_t = \mu + y_{t-1} + e_t$$

- Nonstationary
- Brownian motion: normal innovation
- Random walk

Implication

- conditional and unconditional mean
- conditional and unconditional variance
- h-period ahead forecast

The OLS estimator

$$T(\hat{\gamma}_1 - 1) \xrightarrow{d} \text{ a stable distribution.}$$

but the asymptotic distribution is not normal.

```
library(quantmod, quietly = TRUE)
getSymbols("^GSPC") # S&P 500

y = GSPC$GSPC.Close
plot(y, type = "l")

tail(y)
```

- Null hypothesis: unit root.

$$\Delta y_t = \mu + (\gamma_1 - 1)y_{t-1} + e_t = \mu + \theta y_{t-1} + e_t$$

where $\theta = \gamma_1 - 1$. Under the null, $\theta = 0$.

- The t-statistic is the test statistic for the Dicky-Fuller test.
- Under the null, the t-statistic asymptotically follows a pivotal distribution.

- In this numerical example, the test does not reject the null.

Notice: the test is one-sided.

```r
library(urca) # package for unit root and cointegration
print( summary(ur.df(y, type = "drift", lags = 0) ) ) # y here is
the S&P 500 index
# the test does not reject the null of "unit root"
# loosely speaking, it is evidence in support of random walk
```

```r
y = arima.sim( model = list(ar =.8), n = 100)
plot(y)
```

```r
library(urca); summary( ur.df( y, type = "none", selectlags =
"AIC" ) ) # ,
```

```r
library(dynlm)

DF.sim = function(ar){
  Rep = 500
  n = 100

  B.hat = rep(0, Rep)

  for (r in 1:Rep){
    if (ar < 1) {
      y = arima.sim( model = list(ar = ar), n = n)
    } else if (ar == 1){
      y = ts( cumsum( rnorm(n) ) )
    }
    reg.dyn = dynlm( y ~ L(y,1) )
    B.hat[r] = coef(reg.dyn)[2]
  }
  return(B.hat)
  print("simulation is done with ar = ", ar, "\n")
}


B = DF.sim(1)
plot(density(B), col = "red", xlim = c(0, 1.1))

B = DF.sim(0.5)
```

```
lines(density(B), col = "blue")

B = DF.sim(0.9)
lines( density(B), col = "purple" )
```

Specification of DF test

- The error term must be a white noise for the DF distribution
- DF test's critical values vary with the specfication of drift and/or trend
- Augmented Dicky-Fuller test: add more differenced lag terms

Other tests

- Phillips-Perron test
- KPSS test

Time Series Filtering

Hodrick-Prescott filter: Decompose a time series into *trend* and *cycle*

$$\hat{f}_t = \arg\min_{f_t} \left\{ \sum_{t=1}^{n} (y_t - f_t)^2 + \lambda \sum_{t=3}^{n} \left(\Delta^2 f_t\right)^2 \right\}.$$

Hodrick Prescott (1980, 1997) suggest $\lambda = 1600$ for quarterly data. $\lambda = 1600$ is also the base of adjustment for different time series data frequency.

```
library(mFilter)
data(unemp)

unemp.hp <- hpfilter(unemp)
plot(unemp.hp)
```

Boosted HP Filter

- Phillips and Shi (2019): "[Boosting: Why You Can Use the Hodrick-Prescott Filter](#)," arXiv: 1905.00175
- Chen and Shi (2019): R package **BoostedHP**

Cointegration

In a regression

$$y_t = \beta x_t + e_t$$

- If y_t and x_t are I(1) series
- But a linear combination $e_t = y_t - \beta x_t$ is I(0)

then we say y_t and x_t are cointegrated.

```
[ ]    data(denmark)
```

period: Time index from 1974:Q1 until 1987:Q3.

- `LRM` Logarithm of real money, M2.
- `LRY` Logarithm of real income.
- `LPY` Logarithm of price deflator.
- `IBO` Bond rate.
- `IDE` Bank deposit rate.

```
[ ]    sjd <- denmark[, c("LRM", "LRY", "IBO", "IDE")]
       matplot(sjd[,c(3,4)], type = "l") # Bond rate vs Bank deposit
       rate
```

```
[ ]    y = sjd[,3]; x = sjd[,4]
       reg = lm(y~x)
       plot(residuals(reg), type = "l")
       abline( h = 0, lty = 2)
```

Source of Cointegration

Common shock is the source of cointegration

For example, if $y_{1t} = \mu_1 + \beta_1 t + e_{1t}$ and $y_{2t} = \mu_2 + \beta_2 t + e_{2t}$, where e_{1t} and e_{2t} are two white noises, then the cointegration vector must be $(1, \theta)$ where

$$\theta = -\beta_1/\beta_2.$$

The first coefficient 1 in this cointegration vector is for normalization.

In this example, the common trend is a determistic one. In other examples, they can also share a stochastic trend.

Cointegration

More generally, for an m-vector y_t is cointegrated if there exists a parameter vector γ (normalize the first element to be 1) such that $y_t'\gamma$ is I(0).

- The number of linear independent cointegrated vectors is called the **cointegration rank**.
- The cointegration rank arranges from 1 to $m - 1$.

Error Correction Model

Cliver Granger (Nobel prize 2001)

Consider an ARDL(1,1) model

$$y_t = \mu + \beta_0 x_t + \beta_1 x_{t-1} + \gamma_1 y_{t-1} + e_t.$$

If $\beta_0 = \beta_1 = 0$, no *Granger causality* between X and Y. When X and Y are both nonstationary, standard OLS inference is invalid.

Subtract y_{t-1} from both sides of

$$\begin{aligned}
\Delta y_t &= \mu + \beta_0 x_t + \beta_1 x_{t-1} + (\gamma_1 - 1)y_{t-1} + e_t \\
&= \mu + \beta_0 \Delta x_t + (\beta_1 + \beta_0)x_{t-1} + (\gamma_1 - 1)y_{t-1} + e_t \\
&= \mu + \beta_0 \Delta x_t + (\gamma_1 - 1)(y_{t-1} - \theta x_{t-1}) + e_t
\end{aligned}$$

where $\theta = (\beta_1 + \beta_0)/(1 - \gamma_1)$.

- A short-run relationship $\Delta y_t \sim \mu + \beta_0 \Delta x_t + e_t$.
- An long-run equilibrium error $(\gamma_1 - 1)(y_{t-1} - \theta x_{t-1})$.

When y_t is nonstationary

- First difference recovers stationarity
- Useful to identify spurious regression
- Can be estimated either by OLS or by NLS

Predictive Regression

In the regression

$$y_t = \mu_y + \beta_1 x_{t-1} + e_{yt}$$

- y_t is stationary
- The predictor x_t is highly persistent:

$$x_t = \mu_x + \gamma x_{t-1} + e_{xt}$$

  with γ is close to 1.

- Even if $E[e_{yt}|x_{t-1}] = 0$, OLS estimator of β_1 is biased in finite sample when e_{yt} and e_{xt} are correlated (Stambaugh, 1999).

- Lee, Shi and Gao (2018): "On LASSO for Predictive Regression," arXiv: 1810.03140
- Find new behavior of popular machine learning methods in predictive regression.

Vector Autoregression (VAR)

Christopher Sims (Nobel Prize 2011)

An m-equation system

$$y_t = \mu + \Gamma_1 y_{t-1} + \cdots + \Gamma_p y_{t-p} + v_t$$

where $E[v_t v_t'] = \Omega$.

For prediction purpose, as a reduced-form of structural simultaneous equations.

Estimation

- For consistency and asymptotic normality, use OLS equation by equation

- For asymptotic efficient, use multiple-equation GLS

Invertibility

Write the VAR(p) as

$$(I_m - \Gamma(L))y_t = \mu + v_t$$

where $\Gamma(z) = \Gamma_1 z + \cdots + \Gamma_p z^p$.

Stable means that all roots of the pth order polynomial equation

$$I_m - \Gamma(z) = 0_m$$

lies out of the unit circle.

Impulse Response Function

IRF characterizes the diffusion of an exogenous shock with the dynamic system.

$$y_t = (I_m - \Gamma(L))^{-1}(\mu + v_t)$$
$$= \bar{y} + \left(v_t + \sum_{i=1}^{\infty} A_i v_{t-i}\right)$$

```
library(tsDyn)
data(barry)
plot(barry)
```

```
## For VAR
mod_var <- lineVar(barry, lag = 2)
print(mod_var)
```

```
irf_var = irf(mod_var, impulse = "dolcan", response = c("dolcan",
"cpiUSA", "cpiCAN"), boot = FALSE)
# no matter stationary or not, OLS point estimator is always
valid
print(irf_var)
```

```
[ ]    plot(irf_var)
```

Structural VAR

- Unrestricted VAR: too many parameters? $m + p \cdot m^2 + m(m+1)/2$
- Use economic theory to reduce the number of unknown parameters

VECM Representation

Suppose there are r cointegration relationship in y_t. For the m-equation VAR system

$$y_t = \Gamma y_{t-1} + e_t,$$

we can rewrite it as

$$\Delta y_t = (\Gamma - I_m)y_{t-1} + e_t = \Pi y_{t-1} + e_t.$$

- Since LHS is stationary, the $m \times m$ matrix $\Pi = \Gamma - I_m$ on the RHS must only have rank at most r.
- Otherwise, the RHS will be I(1) and the two sides of the equation are unbalanced.

```
[ ]    ## For VECM
       mod_VECM <- VECM(barry, lag = 2, estim="ML", r=2)
       # VECM is better choice for standard inference when nonstationary
       time series is involved.
       print(mod_VECM)
       irf_vecm = irf(mod_VECM, impulse = "dolcan", response =
       c("dolcan", "cpiUSA", "cpiCAN"), boot = FALSE)
       print(irf_vecm)
```

```
[ ]    plot(irf_vecm)
```

Johansen Test (not covered in 2019)

- VECM is the base for the cointegration rank test (Johansen, 1992).

- Numerical examaple: The result shows that there is only 1 cointegration relationship among the 4 time series.

```
sjd.vecm = ca.jo(sjd, ecdet = "const", type="eigen", K=2,
spec="longrun")
summary(sjd.vecm)
# the result rejects "r = 0", but do not reject "r<=1,2,3"
```

Future of Time Series Study

- Classical methods
- Time series model for discrete choice model
- Time series dimension of big data
    - Unstructured data
    - Panel data